

UNIX and C Programming (COMP1000)

Lecture 3: Pointers

Updated: 13th August, 2019

Department of Computing
Curtin University

Copyright © 2019, Curtin University
CRICOS Provide Code: 00301J

Textbook Reading (Hanly and Koffman)

For more information, see the weekly reading list on Blackboard.

- ▶ **Chapter 6: Pointers and Modular Programming**

Section 6.1 *also* introduces “Pointers to Files”, which you can safely ignore until lecture 5.

- ▶ **Appendix A: More About Pointers**

This is important material, even though it’s in an appendix.

For Test 1, revise everything up to, and including, next week’s lecture.

Outline

Memory

Pointer Introduction

Using Pointers


Heap Allocation

Valgrind


Memory

- ▶ Essentially a single giant sequence of bytes.
- ▶ Each byte has a unique, sequential “address”.
- ▶ Carved up in complex ways, into many parts.

Memory Structure



Code Segment




- ▶ Read-only.
- ▶ Consists of a sequence of functions.
- ▶ Each function is a sequence of machine code instructions.
- ▶ Each instruction occupies one or more bytes.
- ▶ The last instruction is the “return” instruction.

Global Storage

- ▶ Contains statically-allocated data, including:
 - ▶ Global variables.
 - ▶ Local static variables.
 - ▶ String literals.
- ▶ Writable (not read-only), but has a fixed size

Global Storage — Example



```
#include <stdio.h>


double globalVar = 1.5;

void printNext(void)
{
    static int count = 0;
    count++;
    printf("%d\\n", count);
}
```


The Stack

- ▶ “Last in, first out”
- ▶ Used to store:
 - ▶ Temporary/intermediate calculation results
 - ▶ Local variables (except static variables)
 - ▶ Function parameters (sometimes)
 - ▶ Function return location (i.e. where to jump back to when a function finishes)
- ▶ Grows and shrinks over time
- ▶ Grows when a function is called (when local variables are allocated)
- ▶ Shrinks when a function returns (when local variables are destroyed)

The Stack — Example



```
int a = 10;
int b = 12;
...
f(a, b);
...

void f(int x, int y)
{
    int z = x / y;
    printf("%d\n", z);
}
```

The Heap


- ▶ A giant pool of dynamically-allocated memory.
- ▶ No particular fixed size or structure.
- ▶ Allocated and freed as needed by the program.
- ▶ Discussed later in this lecture.

The Registers

- ▶ Not part of main memory; found inside the CPU itself.
- ▶ Extremely small amount of storage — only a few bytes.
- ▶ Very fast.
- ▶ Used to store (depending on the compiler and if there's enough room):
 - ▶ Function parameters.
 - ▶ Temporary/intermediate calculation results.
 - ▶ Variables declared with the register storage class.
- ▶ Registers also keep track of:
 - ▶ The location of the stack.
 - ▶ The currently-executing instruction.
- ▶ For most programming purposes, we can usually ignore the registers.

Pointers

- ▶ Each byte in memory has a unique, sequential “address”.
- ▶ These addresses themselves are often stored in memory.
- ▶ This is called a “pointer”.
- ▶ Where a value occupies multiple bytes (as most do), a pointer will point to the first byte (i.e. lowest memory address).



Pointer Data Types

- ▶ Pointers are a data type, or rather a *set of data types*.
- ▶ To declare a pointer, we must also declare what sort of data it points to.

```
int* size;  
double* speed;  
char* letter;
```

- ▶ A pointer to an `int` and a pointer to a `float` are different data types.

Pointer Declaration Syntax

data-type * name;

- ▶ Declares a variable (“**name**”), storing a memory address.
- ▶ That memory address itself stores a value of type “**data-type**” (e.g. `int`, `float`, etc.).
- ▶ Spacing is meaningless. These are all equivalent:

```
int*x;  
int *x;  /* Widely used. */  
int* x;  /* Used by me (to avoid confusion). */  
int * x;
```

- ▶ Later we'll see “`*x`” used like a variable, but be careful:
 - ▶ The real variable is still `x`, not `*x`.
 - ▶ The `*` has *another* meaning: “dereference”.

Multiple Declarations

- ▶ You can declare multiple pointers on the same line:

```
float *ptr1, *ptr2, *ptr3;
```

- ▶ Notice that they each require a *.
- ▶ The * is part of the data type, but attaches to one name only.
- ▶ Consider this:

```
float* ptr, number;
```

ptr is a pointer, but number is just a float.

NULL

- ▶ A special pointer value — a “pointer to nothing”.
- ▶ Useful for initialising pointers.
- ▶ Actually a preprocessor macro, defined in all the standard header files.
- ▶ Like Java’s “null”. *Not* like SQL’s “NULL”.

Extra Note

- ▶ NULL is *often* defined to be zero (on many platforms).
- ▶ Recall that FALSE is also zero (from the previous lecture).
- ▶ This fact is often used (or misused) inside if, while and do-while statements.

Dereferencing

The Uses of *


1. Multiplication (obviously);
2. Declaring a pointer (e.g. `int*` ptr); and
3. “Dereferencing” (following) a pointer.

The Dereference Operator

- ▶ After you’ve created a pointer, you can use it to access the memory location it points to (“dereferencing”).
- ▶ Place `*` in front of a pointer variable; e.g. “`*ptr`”.
- ▶ “`*ptr`” is what ptr points to. If ptr points to an int, then `*ptr` is an int.
- ▶ You can use `*ptr` just like an ordinary variable.
 - ▶ But remember that ptr is the real variable.

Dereferencing Example

```
int* pointer;  
int number;  
  
...  
/* Assume 'pointer'  
points to a valid  
memory address! */  
  
*pointer = 5;  
number = *pointer;
```



- ▶ 5 is stored at the memory address contained in pointer.
- ▶ number is assigned the same value.

Declaration or Dereference?

- ▶ You can declare and initialise a pointer on one line:

```
int* x = y;
```

OR

```
int *x = y; /* Remember: spacing is meaningless.*/
```

- ▶ Be careful – this does *not* dereference x!
- ▶ It *creates* x, with the datatype int* (a pointer to an integer).
- ▶ It also initialises x with the value of y.
 - ▶ y must be another pointer of the same type (int*).
- ▶ The following is equivalent to the above:


```
int* x;  
x = y;
```

Obtaining Pointers — The Address-Of Operator

- ▶ How do you make a pointer point to something?
- ▶ Use the `&` (address-of) operator to get the location of a variable.
- ▶ Place a `&` in front of a variable name; e.g. `&var`.
- ▶ `&var` is a pointer to `var`.

Example

```
int var = 15;  
int* ptr = &var;
```



Now, `*ptr` is `var`!

Referencing and Dereferencing


- ▶ The `*` and `&` operators are opposites.
- ▶ `&` gets an address, while `*` follows an address to get a value.
- ▶ `*&var` is equivalent to just `var`.
- ▶ `&*pointer` is equivalent to `pointer`.

Order of Operators

- ▶ Consider this: `*ptr++;`
- ▶ Dereference 1st and increment 2nd?
- ▶ Or increment 1st and dereference 2nd? (Correct!)
- ▶ To avoid ambiguity/confusion:
 - ▶ Use brackets: `(*ptr)++;`
 - ▶ Or use pre-increment: `++*ptr;`

Pointer Example

```
int var = 10;  
int* ptr;  
  
ptr = &var; /* Make ptr point to var */  
*ptr = 20; /* Set var to 20 */  
  
printf("%d\n", var); /* Prints "20". */
```



- ▶ The line “`ptr = &var`” stores the address of `var` in `ptr`.
- ▶ Thus, `*ptr` is `var`.
- ▶ Thus, “`*ptr = 20`” actually modifies `var`.

Another Pointer Example

```
int b = 15;
int c = 30;
int* x;
int* y;

x = &b;    /* Make x point to b */
y = x;     /* Make y point to whatever x points to */
x = &c;     /* Make x point to c */

printf("%d\n", *x);    /* Prints "30". */
printf("%d\n", *y);    /* Prints "15". */
```

At the end:

- ▶ x points to c
- ▶ y points to b

L-values

- ▶ Consider the humble assignment statement:

```
x = y + 5;    /* Changes x to be equal to y + 5. */
```

- ▶ The right side can be a complex expression, whereas the left side is often just a variable name. . .
- ▶ . . . but not always!
- ▶ Expressions are allowed on the left, *if* they are “L-values”.
- ▶ An L-value is a value *plus* a memory location that holds it.
 - ▶ e.g. x and y. Variables have values plus memory locations.
- ▶ Non-L-values include:
 - ▶ 5 – literal numbers have no memory location.
 - ▶ y + 5 – arithmetic has no memory location.
 - ▶ (The compiler could choose to store these in memory, but only temporarily, and it may not need to at all.)

L-values and Pointers

- ▶ The dereference (*) operator creates an L-value.

```
int num = 42;  
int* ptr = &num;  
...  
*ptr = 64;
```

- ▶ “*ptr” is an expression, not a variable.
 - ▶ But it’s an L-value, so it can appear on the left.
 - ▶ *ptr has a value (42) *and* a memory address (that of num).
- ▶ The address-of (&) operator *requires* an L-value.

```
int num = 42;  
int* ptr;  
...  
ptr = &num;
```

- ▶ You cannot write &42 or &(num + 5). They have no meaning.

Passing by Value

- ▶ Function parameters are usually “passed by value”.
- ▶ A function’s parameters are *copies* of the values passed into it.
- ▶ The function can overwrite the copies without affecting the originals.

Example

```
double invert(double x) {  
    x = 1.0 / x;  
    return x;  
}  
...  
double a = 0.5, b;  
b = invert(a);
```

Afterwards, `a == 0.5` and `b == 2.0`.

Passing by Reference

- ▶ Using pointers, parameters can be “passed by reference”.
- ▶ Instead of supplying a copy of the value, you supply a pointer to the original.
- ▶ The function can then change the original variable.

Example

```
void invert(double* x) {  
    *x = 1.0 / *x;  
}  
  
...  
double a = 0.5;  
invert(&a);
```

The value of “a” changes from 0.5 to 2.0.

Passing by Reference (2)

- ▶ Used when the value to be passed in occupies many bytes (avoids unnecessary copying).
- ▶ Used when you want the function to export *more than one* value.
- ▶ In C, arrays and strings can only be passed by reference (but we'll get to that later).

Example — `scanf()`


- ▶ Recall that `scanf()` requires a `&` in front of its parameters.
- ▶ These parameters are exports, not imports.
- ▶ `scanf()` needs memory addresses to store values.

Pointers to Pointers

- ▶ A pointer can point to any data type, including another pointer.
- ▶ This is called a “pointer to a pointer (to another datatype)”.
- ▶ (a “pointer to a pointer to a pointer...”, etc. are also possible.)
- ▶ We will learn more about their uses later in the unit.

Example

```
int num = 60;  
int* s;    /* Pointer to an integer */  
int** d;   /* Pointer to a pointer */  
  
d = &s;     /* Make d point to s */  
*d = &num;  /* Make s point to num */  
**d = 30;   /* Assign 30 to num */
```



Void Pointers

- ▶ A pointer to “void”; a generic pointer.
- ▶ Contains a valid memory address, but the type of data stored there is unspecified.
- ▶ Cannot be dereferenced (directly):

```
void assign(int* intPtr, void* voidPtr)
{
    *intPtr = 42;    /* Perfectly fine. */
    *voidPtr = 42; /* Illegal. */
}
```

- ▶ We can't write `*voidPtr`, because we don't know the datatype.
- ▶ To use `void*`, it must be *typecast* to something else.

Void Pointers and Typecasting

- ▶ Pointers can be typecast to different kinds of pointers:

```
void assign(void* voidPtr) {  
    int* intPtr;  
    intPtr = (int*)voidPtr; /* Typecast & assign*/  
    *intPtr = 42;  
}
```

- ▶ Here, we copy voidPtr to intPtr.
 - ▶ The typecast is highlighted.
 - ▶ Then we can access the memory.
 - ▶ We *assume* voidPtr really points to an int (exercise caution).
- ▶ We don't even need intPtr, actually:

```
void assign(void* voidPtr) {  
    *(int*)voidPtr = 42; /* Same as above. */  
}
```


Void Pointers — Example

- ▶ Void pointers are sometimes used to handle multiple datatypes:

```
void printN(char format, void* n) {  
    if(format == 'i') {  
        /* Assume n really points to an int. */  
        printf("%d\n", *(int*)n);  
    }  
    else if(format == 'f') {  
        /* Assume n really points to a float. */  
        printf("%f\n", *(float*)n);  
    }  
}
```

- ▶ Here, printN() can be used to print ints and floats.
- ▶ The format parameter tells us how to treat the void pointer.

Invalid Pointers


- ▶ C offers no guarantee that a pointer points to anything meaningful.
- ▶ It's your job to make sure it does.
- ▶ Newer languages (like Java) use pointers “behind the scenes”, and so protect you from misusing them.
- ▶ C offers no protection.

Example


Say your function returns a pointer to one of its own local variables:

- ▶ The local variable will be destroyed when the function ends.
- ▶ The pointer will point to memory no longer in use.

Invalid Pointers — Example



```
int* ptr;  
ptr = func();  
...  
  
int* func(void)  
{  
    int var = 33;  
    return &var;  
}
```



Heap Memory

- ▶ So far we've dealt with the stack and global storage.
- ▶ The “heap” is a much larger (and more flexible) pool of memory.
- ▶ Your functions can return pointers to heap-allocated memory.
 - ▶ You don't have to allocate and use the memory in the same function.
 - ▶ Can't be done with stack-based variables.
- ▶ Mostly useful for arrays (lecture 4) and structs (lecture 6).

Allocation

- ▶ All memory must be “allocated” before it can be used.
- ▶ Stack memory is allocated simply by declaring a local variable.
- ▶ Heap memory is allocated by calling the `malloc()` function.
- ▶ The OS finds a block of memory of a suitable size, and grants your program permission to access it.
- ▶ This memory *stays* allocated until you “free” it — it doesn't disappear when a function ends.

Deallocation / Freeing

- ▶ Newer languages (like Java) have “garbage collection”.
 - ▶ When a block of heap memory no longer has any pointers to it, it is automatically freed.
- ▶ **C does not** have garbage collection!
- ▶ You must explicitly deallocate heap memory when you're done with it.

The Operating System

- ▶ When a program ends, all its allocated memory *is* forceably freed by the OS.
- ▶ For large or long-running programs, this doesn't help you.
- ▶ Don't rely on it!

Memory Leaks

- ▶ If you forget to deallocate heap memory in C, you get a “memory leak”.
- ▶ This is often an invisible error.
- ▶ Your program may appear to work perfectly.
- ▶ However, it may consume much more memory than needed (and you may run out of memory).
- ▶ There are tools — like `valgrind` — to help detect memory leaks.

Heap Allocation in C

- ▶ Heap allocation (and de-allocation) is done through specific functions.
- ▶ Mainly uses the `stdlib.h` library.
- ▶ Some generic memory-related functions are found in `string.h`.
 - ▶ (These include `memset()` and `memcpy()`, which we'll cover next week.)

The sizeof Operator

- ▶ Reports the size (in bytes) of any data type.
- ▶ A C language construct (*not* a function).
- ▶ On a typical 32-bit machine:

`sizeof(char) == 1`

`sizeof(short) == 2`

`sizeof(int) == 4`

`sizeof(float) == 4`

`sizeof(double) == 8`

`sizeof(void*) == 4`

`sizeof(short**) == 4`

- ▶ These sizes may change across different hardware — that's why `sizeof` is useful!
- ▶ A special unsigned integer type “`size_t`” represents memory sizes:

```
size_t numBytes = sizeof(long);
```

Heap Allocation — malloc()

- ▶ The `malloc()` function allocates a block of heap memory.
- ▶ Takes an `int` parameter — the size in bytes.
- ▶ Returns a `void*` pointing to the newly-allocated memory.
- ▶ Returns `NULL` if the memory could not be allocated.

Usage

Since `malloc()` doesn't know (or care) what you want to store in the memory:

- ▶ Use `sizeof` to determine the number of bytes to allocate.
- ▶ Typecast the returned `void` pointer to the appropriate pointer type.

malloc() Example

To *dynamically* allocate storage for various values:

```
#include <stdlib.h>

...

int* integer = (int*)malloc(sizeof(int));
float* real = (float*)malloc(sizeof(float));
double* bigReal = (double*)malloc(sizeof(double));
```

- ▶ “sizeof(int)” — the number of bytes to allocate.
- ▶ “(int*)” — typecast to an int pointer.
- ▶ Once allocated, you can use *integer, *real and *bigReal just like “ordinary” variables.

Deallocation — `free()`

- ▶ Every block of memory allocated with `malloc()` must eventually be “freed”.
- ▶ The `free()` function takes a pointer to the block, and frees it.
- ▶ Returns nothing.
- ▶ Thereafter, the pointer is invalid.
- ▶ Don't free memory before you're finished with it!

NULL

- ▶ It's good practice to set a freed pointer to `NULL`.
- ▶ You should do this immediately after a call to `free()`.
 - ▶ (Unless the pointer variable itself is about to disappear too.)

free() — Example

```
#include <stdlib.h>

int main(void) {
    double* real = (double*)malloc(sizeof(double));

    ... /* Use the memory */

    free(real);
    real = NULL;

    ... /* Do something else afterwards */
}
```

- ▶ You don't need to tell `free()` how big the block is.
- ▶ Make sure you set the pointer to `NULL` afterwards.

Valgrind

- ▶ It's easy to make mistakes with memory in C:
 - ▶ Using uninitialised values,
 - ▶ Accessing unallocated memory,
 - ▶ Failing to free memory,
 - ▶ Freeing the same memory more than once,
 - ▶ Losing track of allocated memory (memory leaks).
- ▶ The valgrind tool helps you find these.
- ▶ valgrind works with compiled programs.
- ▶ It detects memory errors while a program is running.
- ▶ To use it, type:

```
[user@pc]$ valgrind [options] ./program ...
```

Leave **[options]** blank for a summary of the errors.

Debugging Information

- ▶ To get the most out of valgrind, you need “debugging information”.
- ▶ This is a compile option, which inserts extra information into the executable file.
- ▶ On the command-line:

```
[user@pc]$ gcc -g -c file.c
```
- ▶ In a Makefile:

```
CFLAGS = -Wall -pedantic -ansi -g
```
- ▶ Without this, valgrind can't give you any line numbers.

Valgrind Output

This is valgrind detecting a memory leak:

LEAK SUMMARY:

definitely lost: 10 bytes in 1 blocks

indirectly lost: 0 bytes in 0 blocks

possibly lost: 0 bytes in 0 blocks

still reachable: 0 bytes in 0 blocks

suppressed: 0 bytes in 0 blocks

Rerun with `--leak-check=full` to see details of
leaked memory

Then, to find *where* the leak occurred:

```
[user@pc]$ valgrind --leak-check=full ./program ...
```


Coming Up

- ▶ Next week's lecture will expand on pointers, covering arrays and strings.
- ▶ Once again, make sure you complete the tutorial exercises.